

DATA STRUCTURES

Hash tables: Coalesced chaining, Open addressing

Lect. PhD. Diana-Lucia Miholca

2022 - 2023



Babeş - Bolyai University
Faculty of Mathematics and Computer Science

In previous lecture...

- Hash tables
 - Direct-address table
 - Introduction to hash tables
 - Separate chaining

Today

- Coalesced chaining
- Open addressing



Collision resolution by coalesced chaining: each element is stored inside the table and has associated the index of the *next* element.



Collision resolution by coalesced chaining: each element is stored inside the table and has associated the index of the *next* element.



When **adding a new element** and the index it hashes to is occupied:

- the element is added at any empty index and
- the *next* indexes are set so as to find it starting from its hashing index



Collision resolution by coalesced chaining: each element is stored inside the table and has associated the index of the *next* element.



When **adding a new element** and the index it hashes to is occupied:

- the element is added at any empty index and
- the *next* indexes are set so as to find it starting from its hashing index



Since elements are in the table, α can be at most 1.

Coalesced chaining - example



Consider a hash table that uses coalesced chaining for collision resolution, with:

- $m = 16$
- the division method for hashing



Insert in the hash table, in the given order, the following elements: 76, 12, 109, 43, 22, 18, 55, 81, 91, 27, 13, 16, 39.



What is the hash value for 12? What is the hash value for 16?

Coalesced chaining - example



Consider a hash table that uses coalesced chaining for collision resolution, with:

- $m = 16$
- the division method for hashing



Insert in the hash table, in the given order, the following elements: 76, 12, 109, 43, 22, 18, 55, 81, 91, 27, 13, 16, 39.



What is the hash value for 12? What is the hash value for 16?



Let's compute the hash value for every element (key):

Key	76	12	109	43	22	18	55	81	91	27	13	16	39
Hash	12	12	13	11	6	2	7	1	11	11	13	0	7

Example



Initially, the hash table is empty

- The first empty position is 0 and all *next* indexes are -1.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

firstEmpty = 0



76 is added at index 12



12 should also be added at index 12. Since it is already occupied, we add 12 at index *firstEmpty* (0) and set the *next* of 76 to 0. Then we reset *firstEmpty* to the next empty position.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
12												76			
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	0	-1	-1	-1

firstEmpty = 1

Example

Key	76	12	109	43	22	18	55	81	91	27	13	16	39
Hash	12	12	13	11	6	2	7	1	11	11	13	0	7



We continue similarly, having no collisions up to 81.



We need to update *firstEmpty* every time we occupy it.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
12	81	18				22	55				43	76	109		
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	0	-1	-1	-1

firstEmpty = 3



When adding 91, we add it at index *firstEmpty* and set the *next* link at index 11 to 3.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
12	81	18	91			22	55				43	76	109		
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	3	0	-1	-1	-1

firstEmpty = 4

Example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
12	81	18	91	27	13	22	55	16	39		43	76	109		
8	-1	-1	4	-1	-1	-1	9	-1	-1	-1	3	0	5	-1	-1

firstEmpty = 10

Coalesced chaining - representation

- A hash table with coalesced chaining is represented in the following way:



Representation of a hash table with coalesced chaining:

HashTable:

T: TKey[]

next: Integer[]

m: Integer

firstEmpty: Integer

h: TFunction



For simplicity, in the following, we will consider only the keys.

Coalesced chaining - insert



Adding a key to a hash table with coalesced chaining:

subalgorithm insert (ht, k) **is:**

//pre: ht is a HashTable, k is a TKey

//post: k was added into ht

Coalesced chaining - insert



Adding a key to a hash table with coalesced chaining:

subalgorithm insert (ht, k) **is:**

//pre: ht is a HashTable, k is a TKey

//post: k was added into ht

index \leftarrow ht.h(k)

if ht.T[index] = $NULL_{TKey}$ **then** *//NULL_{TKey} means empty position*

ht.T[index] \leftarrow k

if index = ht.firstEmpty **then**

changeFirstEmpty()

end-if

else

if ht.firstEmpty = ht.m **then**

@resize and rehash

end-if

ht.T[ht.firstEmpty] \leftarrow k

//continued on the next slide...

Coalesced chaining - insert



Adding a key to a hash table with coalesced chaining:

current \leftarrow index

while ht.next[current] \neq -1 **execute**

 current \leftarrow ht.next[current]

end-while

ht.next[current] \leftarrow ht.firstEmpty

ht.next[ht.firstEmpty] \leftarrow - 1

changeFirstEmpty(ht)

end-if

end-subalgorithm



Complexity:

Coalesced chaining - insert



Adding a key to a hash table with coalesced chaining:

current \leftarrow index

while ht.next[current] \neq -1 **execute**

 current \leftarrow ht.next[current]

end-while

ht.next[current] \leftarrow ht.firstEmpty

ht.next[ht.firstEmpty] \leftarrow - 1

changeFirstEmpty(ht)

end-if

end-subalgorithm



Complexity: $\Theta(1)$ on average (under SUH assumption),
 $\Theta(n)$ - in the worst case

Coalesced chaining - ChangeFirstEmpty



Updating the first empty position in a hash table with coalesced chaining:

subalgorithm changeFirstEmpty(ht) **is:**

//pre: ht is a HashTable

//post: the value of ht.firstEmpty is set to the next free position

ht.firstEmpty \leftarrow ht.firstEmpty + 1

while ht.firstEmpty < ht.m **and** ht.T[ht.firstEmpty] \neq $NULL_{TKey}$

execute

ht.firstEmpty \leftarrow ht.firstEmpty + 1

end-while

end-subalgorithm



Complexity:

Coalesced chaining - ChangeFirstEmpty



Updating the first empty position in a hash table with coalesced chaining:

subalgorithm changeFirstEmpty(ht) **is:**

//pre: ht is a HashTable

//post: the value of ht.firstEmpty is set to the next free position

ht.firstEmpty \leftarrow ht.firstEmpty + 1

while ht.firstEmpty < ht.m **and** ht.T[ht.firstEmpty] \neq $NULL_{TKey}$

execute

ht.firstEmpty \leftarrow ht.firstEmpty + 1

end-while

end-subalgorithm



Complexity: $O(n)$

Coalesced chaining

 *Remove* and *search* will be discussed in Seminar 5.



Collision resolution by open addressing: each element is stored inside table and there are no links.



Collision resolution by open addressing: each element is stored inside table and there are no links.



When adding a new element, we will:

- ▶ successively generate candidate positions
- ▶ check (*probe*) their availability and
- ▶ place the element in the first available one

Open addressing



In order to generate multiple positions, the hash function is extended with an additional parameter, i , which is the *probe number* and starts from 0.

$$h : U \times \{0, 1, \dots, m - 1\} \rightarrow \{0, 1, \dots, m - 1\}$$



For an element k , positions from the *probe sequence* $\langle h(k, 0), h(k, 1), h(k, 2), \dots, h(k, m - 1) \rangle$ will be successively examined.



The *probe sequence* should be a permutation of $\{0, \dots, m - 1\}$, so that eventually every slot is probed.

Open addressing - Linear probing



A first scheme for defining the hash function is to use **linear probing**:

$$h(k, i) = (h'(k) + i) \bmod m \quad \forall i = 0, \dots, m - 1$$

- where $h'(k)$ is a *simple* hash function
 - For example: $h'(k) = k \bmod m$



The *probe sequence* for linear probing is:

$$< h'(k), h'(k) + 1, h'(k) + 2, \dots, m - 1, 0, 1, \dots, h'(k) - 1 >$$

Open addressing - Linear probing - example



Consider a hash table that uses open addressing for collision resolution, with:

- $m = 16$
- linear probing with $h'(k) = k \bmod m$



Insert into the table, in the given order, the following elements: 76, 12, 109, 43, 22, 18, 55, 81, 91, 27, 13, 16, 39.



Let's compute the value of the hash function for every element (key) when $i = 0$:

Key	76	12	109	43	22	18	55	81	91	27	13	16	39
Hash	12	12	13	11	6	2	7	1	11	11	13	0	7

Open addressing - Linear probing - example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
27	81	18	13	16		22	55	39			43	76	12	109	91

Open addressing - Linear probing - example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
27	81	18	13	16		22	55	39			43	76	12	109	91



Disadvantages of linear probing:

- *Primary clustering* - long runs of occupied slots
- There are only m distinct probe sequences (once you have the starting position everything is fixed)



Advantages of linear probing:

- Probe sequence is always a permutation
- Can benefit from caching

Open addressing - Quadratic probing



In case of **quadratic probing** the hash function becomes:

$$h(k, i) = (h'(k) + c_1 * i + c_2 * i^2) \bmod m \quad \forall i = 0, \dots, m - 1$$

- where $h'(k)$ is a *simple* hash function
 - for example: $h'(k) = k \bmod m$ and c_1 and c_2 are constants; c_2 should not be 0



Considering a simplified version of $h(k, i)$ with $c_1 = 0$ and $c_2 = 1$ the probe sequence would be:

$$< h'(k), h'(k) + 1, h'(k) + 4, h'(k) + 9, h'(k) + 16, \dots >$$

Open addressing - Quadratic probing



The values of m , c_1 and c_2 should be chosen so that the probe sequence is a permutation.



If m is a prime number only the first half of the probe sequence is unique \Rightarrow once the hash table is half full, there is no guarantee that an empty index will be found.



For example, for $m = 17$, $c_1 = 3$, $c_2 = 1$ and $k = 13$, the probe sequence is

$\langle 13, 0, 6, 14, 7, 2, 16, 15, 16, 2, 7, 14, 6, 0, 13, 11, 11 \rangle$

Open addressing - Quadratic probing



If m is a power of 2 and $c_1 = c_2 = 0.5$, the probe sequence will always be a permutation.



For example for $m = 8$ and $k = 3$:

- $h(3, 0) = (3 \% 8 + 0.5 * 0 + 0.5 * 0^2) \% 8 = 3$
- $h(3, 1) = (3 \% 8 + 0.5 * 1 + 0.5 * 1^2) \% 8 = 4$
- $h(3, 2) = (3 \% 8 + 0.5 * 2 + 0.5 * 2^2) \% 8 = 6$
- $h(3, 3) = (3 \% 8 + 0.5 * 3 + 0.5 * 3^2) \% 8 = 1$
- $h(3, 4) = (3 \% 8 + 0.5 * 4 + 0.5 * 4^2) \% 8 = 5$
- $h(3, 5) = (3 \% 8 + 0.5 * 5 + 0.5 * 5^2) \% 8 = 2$
- $h(3, 6) = (3 \% 8 + 0.5 * 6 + 0.5 * 6^2) \% 8 = 0$
- $h(3, 7) = (3 \% 8 + 0.5 * 7 + 0.5 * 7^2) \% 8 = 7$

Open addressing - Quadratic probing



If m is a prime number of the form $4 * j + 3$, $c_1 = 0$ and $c_2 = (-1)^j$ (so the probe sequence is +0, -1, +4, -9, etc.) the probe sequence is a permutation.



For example for $m = 7$ and $k = 3$:

- $h(3, 0) = (3 \% 7 + 0^2) \% 7 = 3$
- $h(3, 1) = (3 \% 7 - 1^2) \% 7 = 2$
- $h(3, 2) = (3 \% 7 + 2^2) \% 7 = 0$
- $h(3, 3) = (3 \% 7 - 3^2) \% 7 = 1$
- $h(3, 4) = (3 \% 7 + 4^2) \% 7 = 5$
- $h(3, 5) = (3 \% 7 - 5^2) \% 7 = 6$
- $h(3, 6) = (3 \% 7 + 6^2) \% 7 = 4$

Open addressing - Quadratic probing - Example



Consider a hash table that uses open addressing for collision resolution, with:

- $m = 16$
- quadratic probing with $h'(k) = k \bmod m$ and $c_1 = c_2 = 0.5$.



Insert into the table, in the given order, the following elements: 76, 12, 109, 43, 22, 18, 55, 81, 91, 27, 13, 16, 39.

Open addressing - Quadratic probing - example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
13	81	18	16		91	22	55	39		27	43	76	12	109	

Open addressing - Quadratic probing - example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
13	81	18	16		91	22	55	39		27	43	76	12	109	



Disadvantages of quadratic probing:

- *Secondary clustering* - if two elements have the same initial probe positions, their whole probe sequence will be identical:
$$h(k_1, 0) = h(k_2, 0) \Rightarrow h(k_1, i) = h(k_2, i).$$
- There are only m distinct probe sequences
- The performance is sensitive to the values of m , c_1 and c_2 .



Double hashing uses a hash function of form:

$$h(k, i) = (h'(k) + i * h''(k)) \% m \quad \forall i = 0, \dots, m - 1$$

- where $h'(k)$ and $h''(k)$ are *simple* hash functions; $h''(k)$ should never return 0.



For a key, k , the initial probe goes to position $h'(k)$ and the successive probe positions are offset from previous positions by the amount $h''(k)$, modulo m .

Open addressing - Double hashing



Similar to quadratic probing, not every combination of m and h'' will produce a complete permutation.



$h''(k)$ must be relatively prime to m . This can be ensured by:

- Choosing m as a power of 2 and designing h'' in such a way that it always returns an odd number.
- Choosing m as a prime number and designing h'' in such a way that it always returns a value from the $\{1, m-1\}$.

Open addressing - Double hashing

 For example:

$$h'(k) = k \% m$$

$$h''(k) = 1 + k \% (m - 1)$$

- For $m = 11$ and $k = 36$ we have:

$$h'(36) = 3$$

$$h''(36) = 7$$

- The probe sequence is: $\langle 3, 10, 6, 2, 9, 5, 1, 8, 4, 0, 7 \rangle$

Open addressing - Double hashing - example



Consider a hash table that uses open addressing with double hashing for collision resolution, with:

- $m = 17$
- $h'(k) = k \% m$ and $h''(k) = 1 + (k \% 16)$



Insert into the table, in the given order, the following elements: 75, 12, 109, 43, 22, 18, 55, 81, 92, 27, 13, 16, 39.



Values of the two hash functions for each element:

key	75	12	109	43	22	18	55	81	92	27	13	16	39
$h'(key)$	7	12	7	9	5	1	4	13	7	10	13	16	5
$h''(key)$	12	13	14	12	7	3	8	2	13	12	14	1	8

Open addressing - Double hashing - example



The final hash table:

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
16	18		55	109	22		75		43	27	39	12	81		13	92



The main advantage of double hashing is that even if $h(k_1, 0) = h(k_2, 0)$ the probe sequences will be different if $k_1 \neq k_2$.



For example:

- 75: $\langle 7, 2, 14, 9, 4, 16, 11, 6, 1, 13, 8, 3, 15, 10, 5, 0, 12 \rangle$
- 109: $\langle 7, 4, 1, 15, 12, 9, 6, 3, 0, 14, 11, 8, 5, 2, 16, 13, 10 \rangle$



Since for every $(h'(k), h''(k))$ pair we have a separate probe sequence, double hashing generates $\approx m^2$ different permutations.

Open addressing - representation

- A hash table with open addressing for collision resolution is represented in the following way:



Representation of a hash table with open addressing:

HashTable:

T: TKey[]

m: Integer

h: TFunction



For simplicity, in the following, we will consider only the keys.



Adding an element

subalgorithm insert (ht, e) **is:**

//pre: ht is a HashTable, e is a TKey

//post: e was added in ht

Open addressing - insert



Adding an element

subalgorithm insert (ht, e) **is:**

//pre: ht is a HashTable, e is a TKey

//post: e was added in ht

$i \leftarrow 0$

$\text{pos} \leftarrow \text{ht.h}(e, i)$

while $i < \text{ht.m}$ **and** $\text{ht.T}[\text{pos}] \neq \text{NULL}_{\text{TKey}}$ **execute**

//NULL_{TKey} means empty space

$i \leftarrow i + 1$

$\text{pos} \leftarrow \text{ht.h}(e, i)$

end-while

if $i = \text{ht.m}$ **then**

@resize and rehash and compute the position for e again

else

$\text{ht.T}[\text{pos}] \leftarrow e$

end-if

end-subalgorithm



What should the *search* operation do?



How can we *remove* an element from the hash table?



How can we *remove* an element from the hash table?



We cannot just mark the position empty (by storing $NULL_{TKey}$ into it) - *search* might not find other elements



How can we *remove* an element from the hash table?



We cannot just mark the position empty (by storing $NULL_{TKey}$ into it) - *search* might not find other elements



Remove is usually implemented to mark the deleted position with a special value, *DELETED*.



How can we *remove* an element from the hash table?



We cannot just mark the position empty (by storing $NULL_{TKey}$ into it) - *search* might not find other elements



Remove is usually implemented to mark the deleted position with a special value, *DELETED*.



How does the usage of the value *DELETED* affect the implementation of the *insert* and *search* operation?

Open addressing - Performance



Theorem: In a hash table with open addressing with load factor $\alpha = n/m$ ($\alpha < 1$), the *average* number of probes is at most

- for *insert* and *unsuccessful search*

$$\frac{1}{1 - \alpha}$$

- for *successful search*

$$\frac{1}{\alpha} * \ln \frac{1}{1 - \alpha}$$



If α is constant, the average complexity is $\Theta(1)$



Worst case complexity is $\Theta(n)$

Containers represented using hash tables



Hash tables are used for representing the following containers:

- ADT Map (Sorted Map)



Python's dictionaries (`{:}`), Java `HashMap`, `unordered_map` in C++ STL

- ADT MultiMap (Sorted MultiMap)



`HashMultimap` in Guava (Google Core Libraries for Java)
`unordered_multimap` in C++ STL

- ADT Set



`HashSet` in Java Collections API, Python's sets (`{}`)

- ADT Bag



`HashMultiset` in Guava (for Java)

Hash table - Applications



Real-world applications of hash tables:



Programming languages

- Implementation of built-in data types (*dict* in Python, *HashMap* in Java)



Compilers

- For storing the programming language's keywords and for mapping the variables names with memory locations



File system

- For mapping file names to the the file path and to the physical location of that file on the disk



Password Verification:

- For storing hashed passwords



Data Integrity Checks

- To generate checksums on data files



Bibliography

- David M. Mount, *Lecture notes for the course Data Structures (CMSC 420)*, at the Dept. of Computer Science, University of Maryland, College Park, 2001
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein, *Introduction to Algorithms*, Third Edition, The MIT Press, 2009
- Narasimha Karumanchi, *Data Structures and Algorithms Made Easy: Data Structures and Algorithmic Puzzles*, Fifth Edition, 2016
- Clifford A. Shaffer, *A Practical Introduction to Data Structures and Algorithm Analysis*, Third Edition, 2010

Thank you

▶ CHAINING
THANKS
THAT SHE
YOU COALESCE